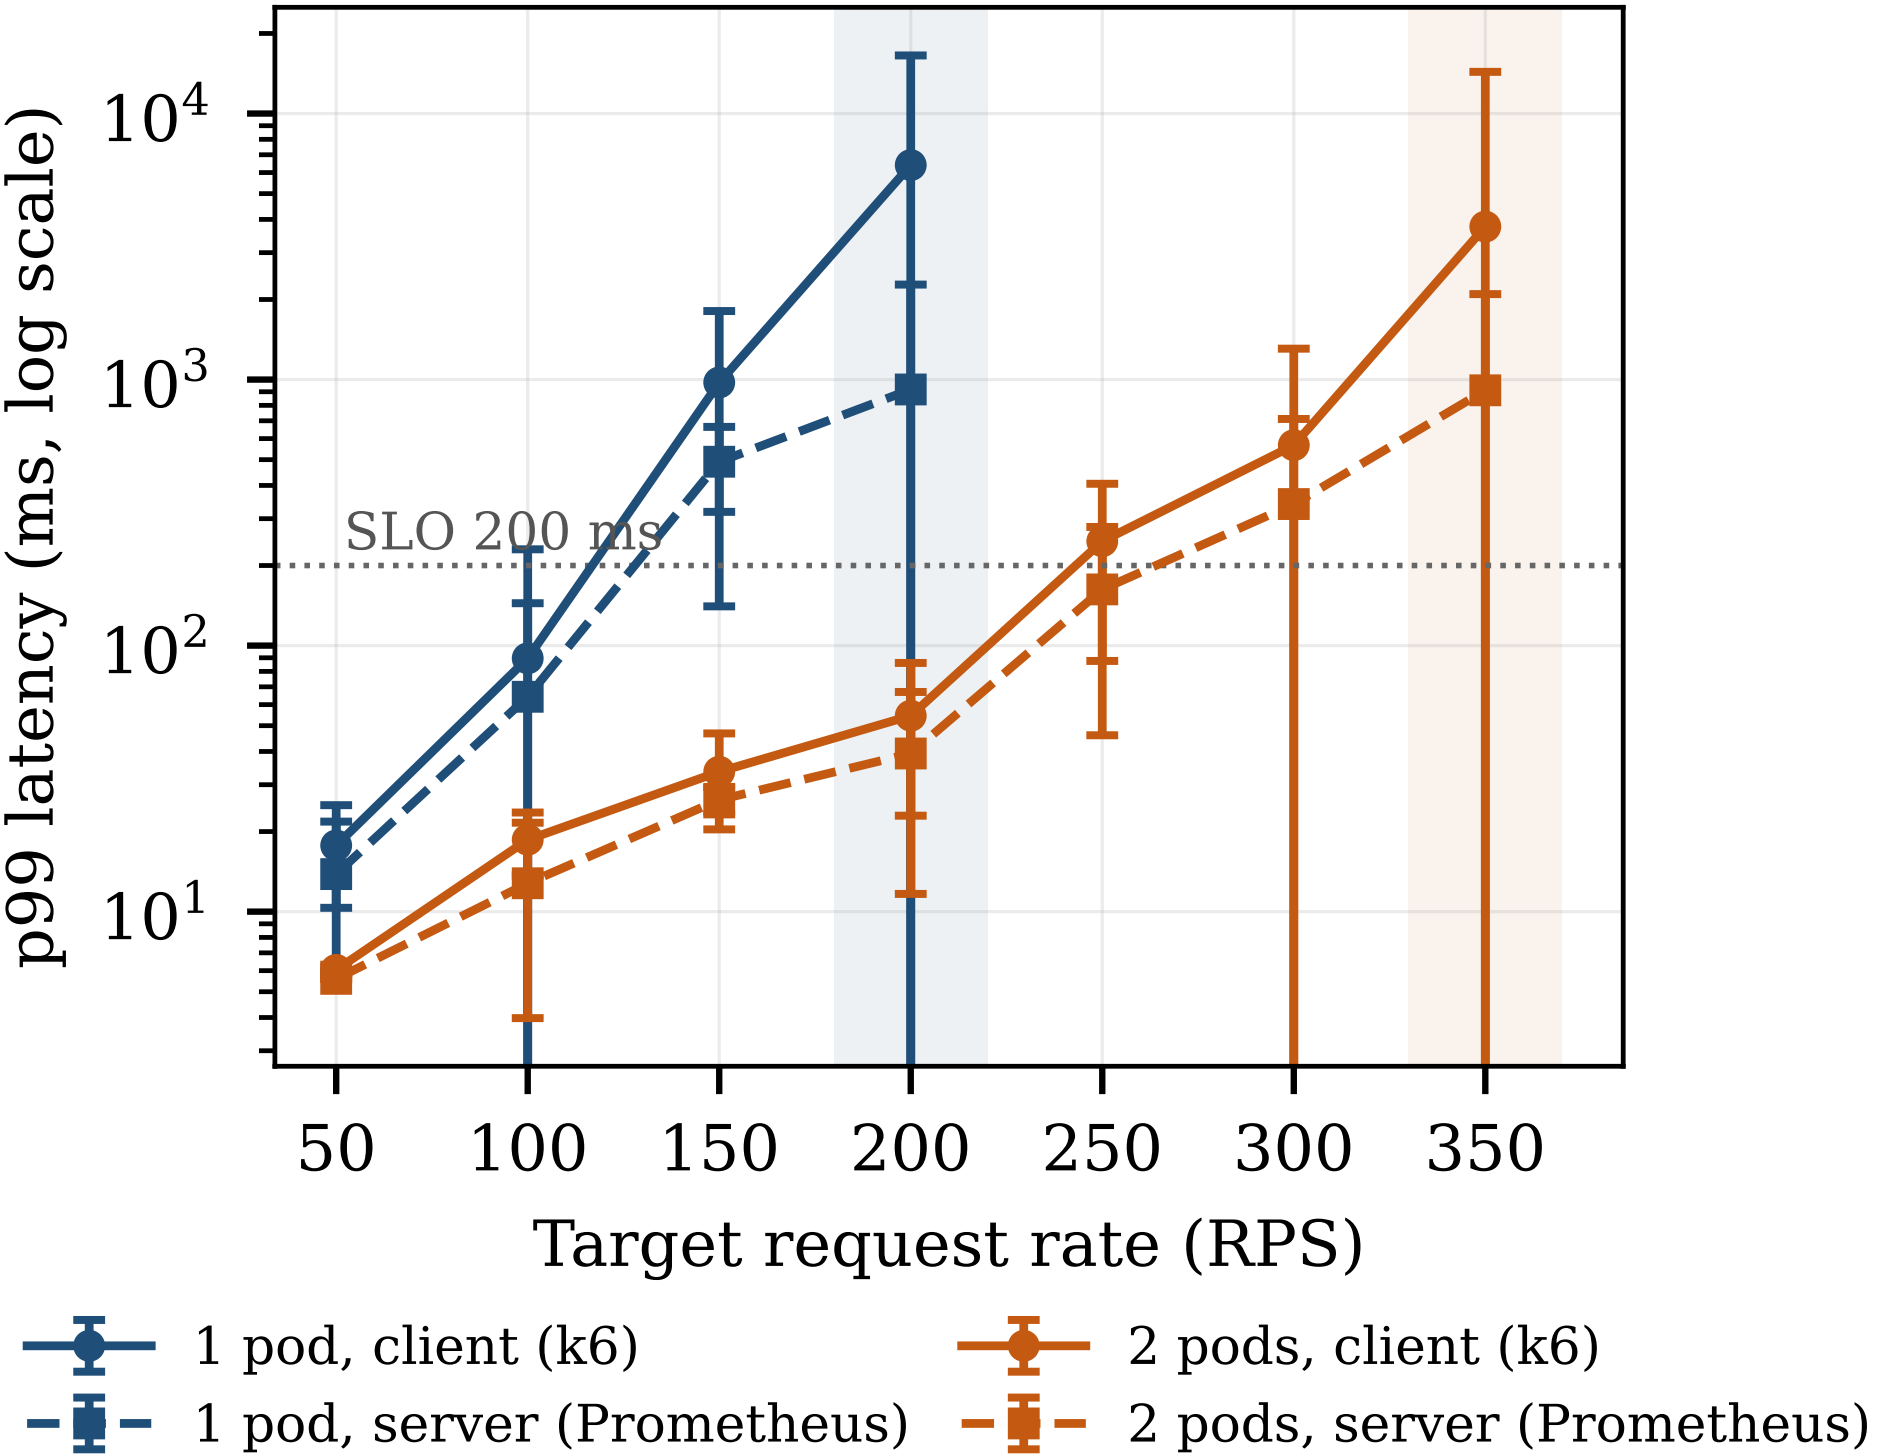
p99 latency (ms, log scale)
10^4
10^3
10^2
10^1
SLO 200 ms
50
100
150
200
250
300
350
Target request rate (RPS)
1 pod, client (k6)
1 pod, server (Prometheus)
2 pods, client (k6)
2 pods, server (Prometheus)